



What is Machine Learning?



Machine Learning

- Before we jump into Neural Networks, Tensorflow, Keras API etc... its a good idea to understand a few fundamental ideas regarding machine learning.
- In this section we'll cover some important theory and concepts surrounding machine learning.



Machine Learning

- Section Overview:
 - What is Machine Learning?
 - What is Deep Learning?
 - Difference between Supervised and Unsupervised Learning
 - Supervised Learning Process
 - Evaluating performance
 - Overfitting



What is Machine Learning?

- Machine learning is a method of data analysis that automates analytical model building.
- Using algorithms that iteratively learn from data, machine learning allows computers to find hidden insights without being explicitly programmed where to look.



What is it used for?

- Fraud detection.
- Web search results.
- Real-time ads on web pages
- Credit scoring.
- Prediction of equipment failures.
- New pricing models.
- Network intrusion detection.
- Recommendation Engines
- Customer Segmentation
- Text Sentiment Analysis
- Customer Churn
- Pattern and image recognition.
- Email spam filtering.



What are Neural Networks?

- Neural Networks are a way of modeling biological neuron systems mathematically.
- These networks can then be used to solve tasks that many other types of algorithms can not (e.g. image classification)
- Deep Learning simply refers to neural networks with more than one hidden layer.



Machine Learning

- There are different types of machine learning we will focus on during the next sections of the course:
 - Supervised Learning
 - Unsupervised Learning



Machine Learning

- Machine Learning
 - Automated analytical models.
- Neural Networks
 - A type of machine learning architecture modeled after biological neurons.
- Deep Learning
 - A neural network with more than one hidden layer.



Machine Learning

- Let's begin by learning about one of the most common machine learning tasks- Supervised Learning!



Supervised Learning



Supervised Learning

- **Supervised learning** algorithms are trained using **labeled** examples, such as an input where the desired output is known.
- For example, a segment of text could have a category label, such as:
 - **Spam** vs. **Legitimate** Email
 - **Positive** vs. **Negative** Movie Review



Supervised Learning

- The network receives a set of inputs along with the corresponding correct outputs, and the algorithm learns by comparing its actual output with correct outputs to find errors.
- It then modifies the model accordingly.




Supervised Learning

- Supervised learning is commonly used in applications where historical data predicts likely future events.



Machine Learning Process





Machine Learning Process

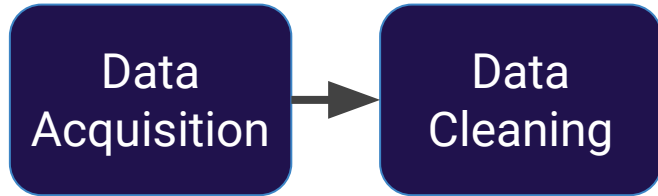
- Get your data! Customers, Sensors, etc...

Data
Acquisition



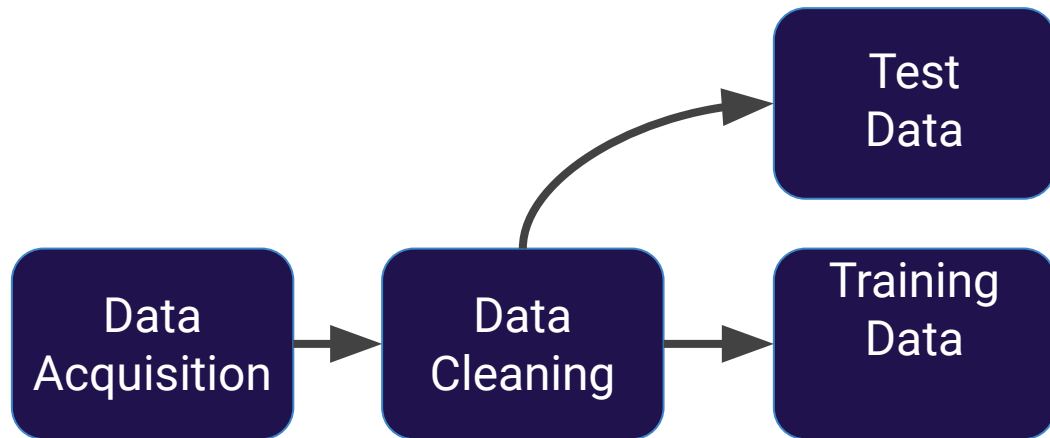
Machine Learning Process

- Clean and format your data (using Pandas)



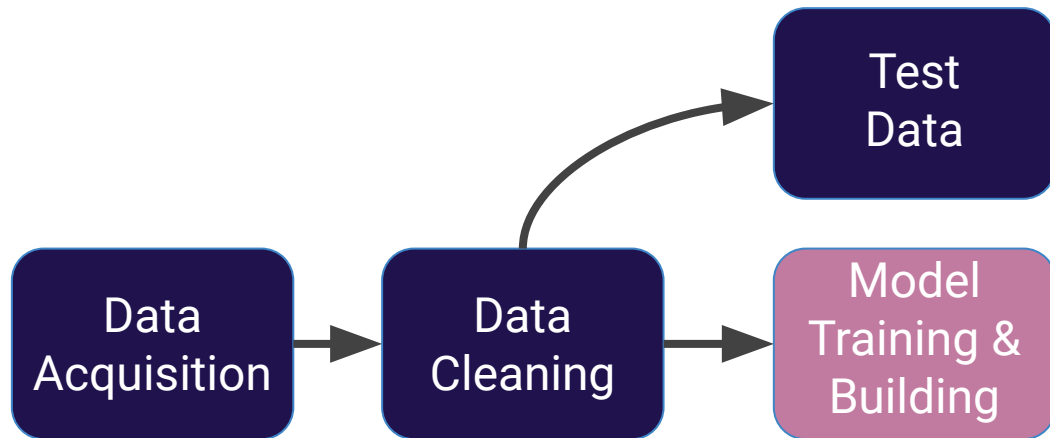


Machine Learning Process



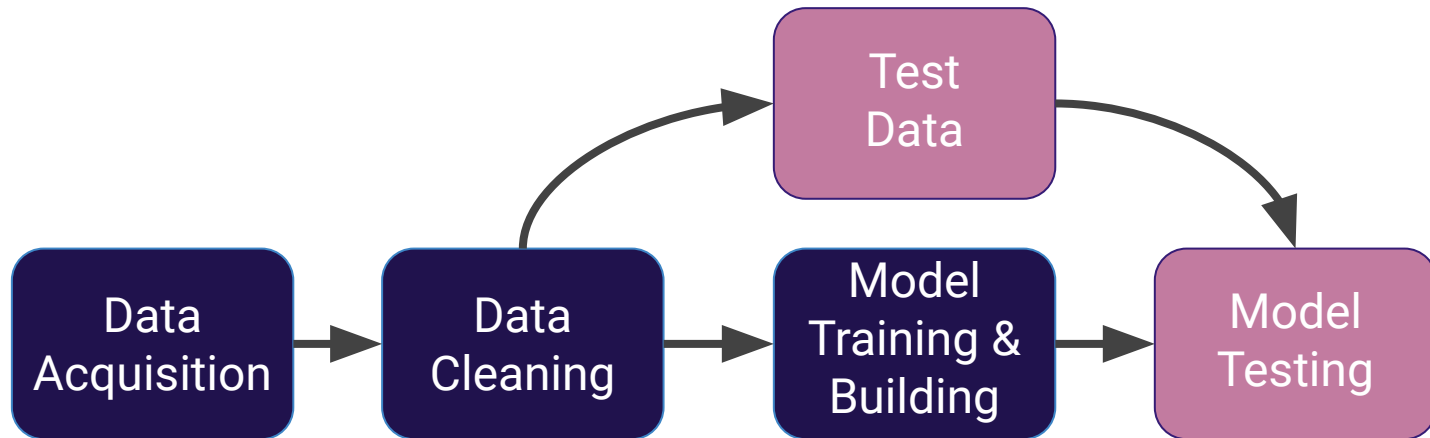


Machine Learning Process



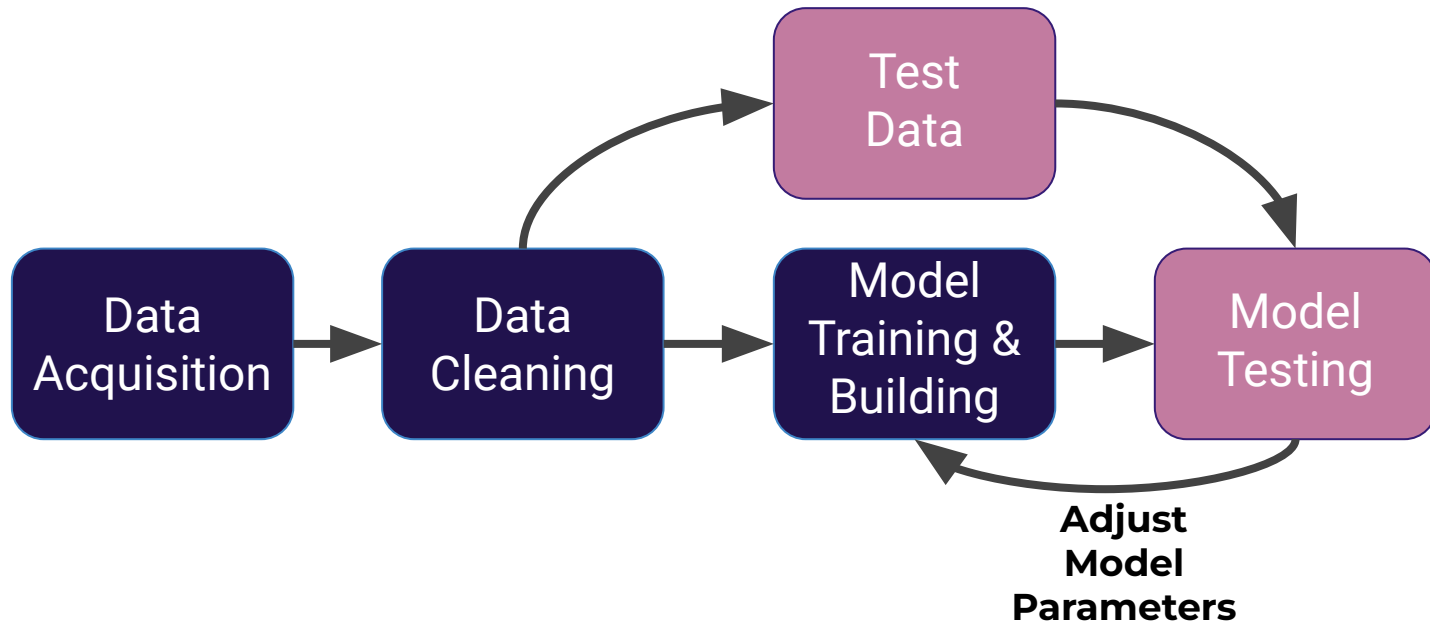


Machine Learning Process






Machine Learning Process





Machine Learning Process





Supervised Learning

- What we just showed is a simplified approach to supervised learning, it contains an issue!
- Is it fair to use our single split of the data to evaluate our models performance?
- After all, we were given the chance to update the model parameters again and again.



Supervised Learning

- To fix this issue, data is often split into **3 sets**
 - Training Data
 - Used to train model parameters
 - Validation Data
 - Used to determine what model hyperparameters to adjust
 - Test Data
 - Used to get some final performance metric



Supervised Learning

- This means after we see the results on the **final test set** we don't get to go back and adjust any model parameters!
- This final measure is what we label the true performance of the model to be.



Supervised Learning

- In this course, in general we will simplify our data by using a simple **train/test split**.
- We will simply train and then evaluate on a test set (leaving the option to students to go back and adjust parameters).
- After going through the course, you will be able to easily perform another split to get **3 data sets** if you desire.



Overfitting and Underfitting



Machine Learning

- Now that we understand the full process for supervised learning, let's touch upon the important topics of **overfitting** and **underfitting**.



Machine Learning

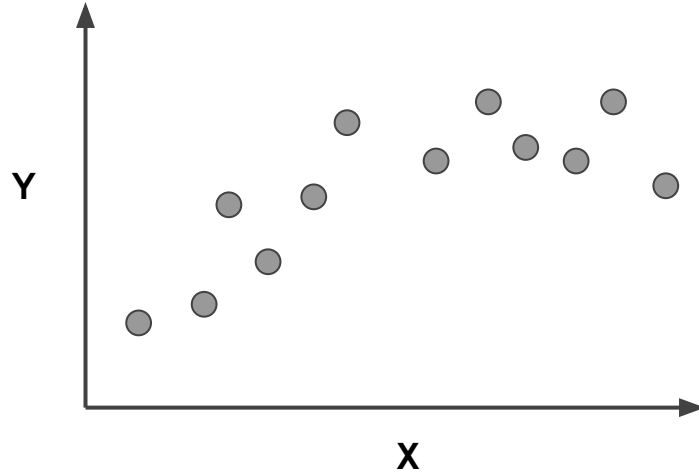
- **Overfitting**

- The model fits too much to the noise from the data.
- This often results in **low error on training sets but high error on test/validation sets.**



Machine Learning


Data





Machine Learning

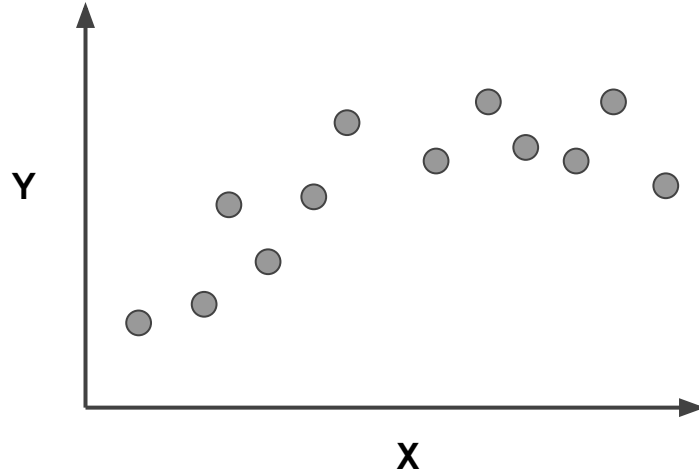
Good Model





Machine Learning


- **Overfitting**





Machine Learning

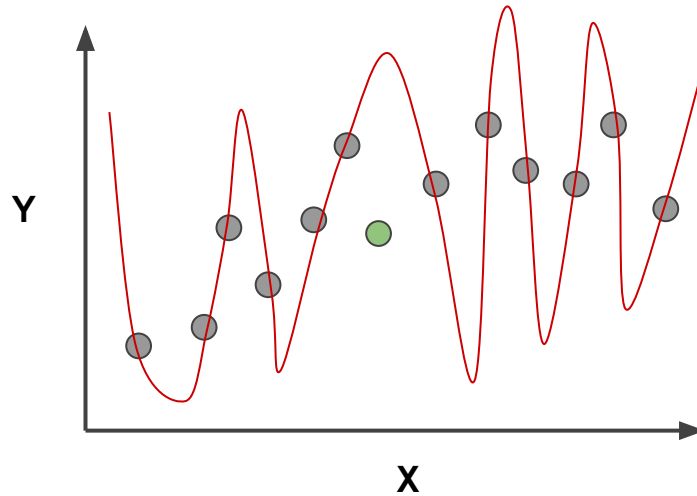
- **Overfitting**





Machine Learning


- **Overfitting**





Machine Learning

- **Overfitting**





Machine Learning

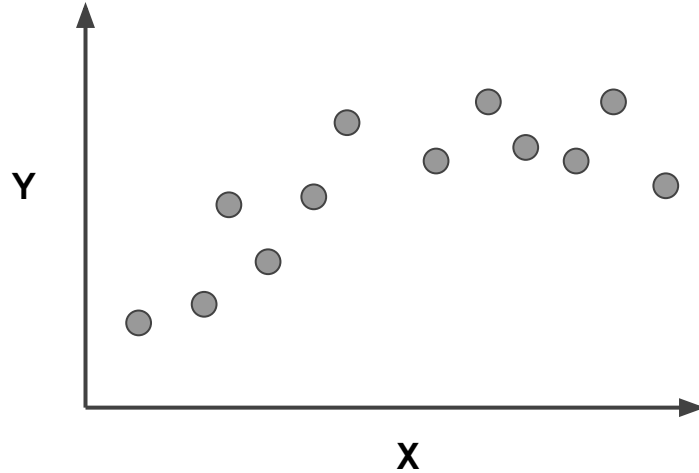
- **Underfitting**

- Model does not capture the underlying trend of the data and does not fit the data well enough.
- Low variance but high bias.
- Underfitting is often a result of an excessively simple model.



Machine Learning

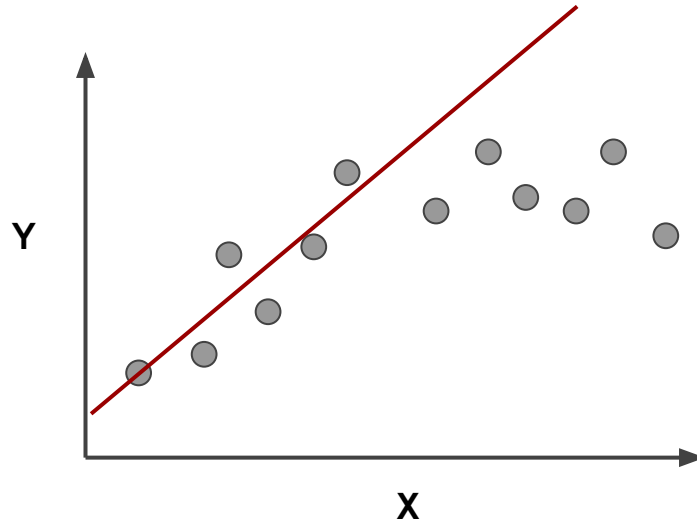
Data





Machine Learning

Underfitting





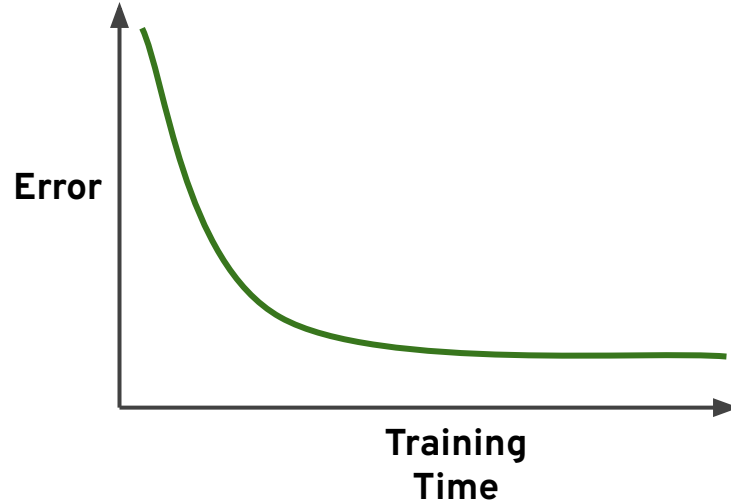
Machine Learning

- This data was easy to visualize, but how can we see underfitting and overfitting when dealing with multi dimensional data sets?
- First let's imagine we trained a model and then measured its error over training time.



Machine Learning


- Good Model





Machine Learning

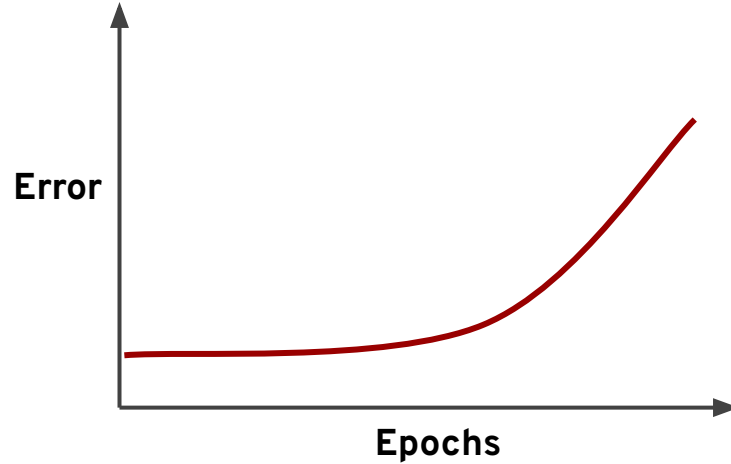
- Good Model





Machine Learning

- Bad Model





Machine Learning

- When thinking about **overfitting** and **underfitting** we want to keep in mind the relationship of model performance on the training set versus the test/validation set.



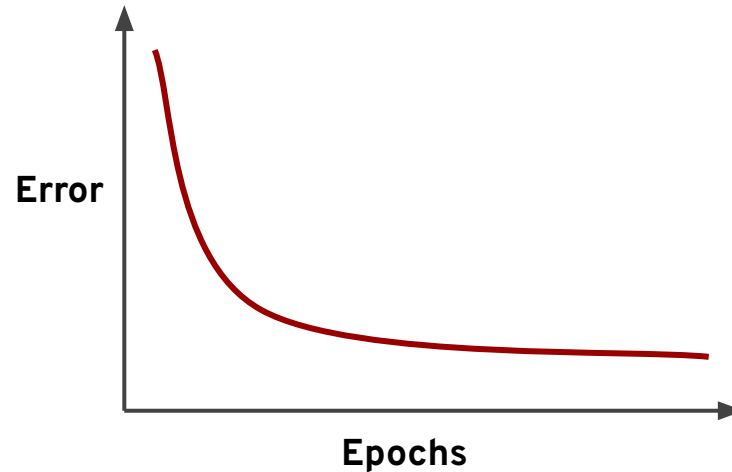
Machine Learning

- Let's imagine we split our data into a **training set** and a **test set**



Machine Learning

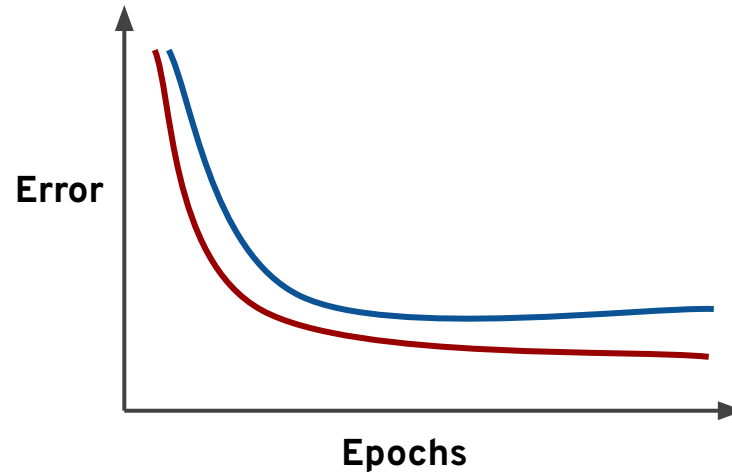
- We first see performance on the **training set**





Machine Learning

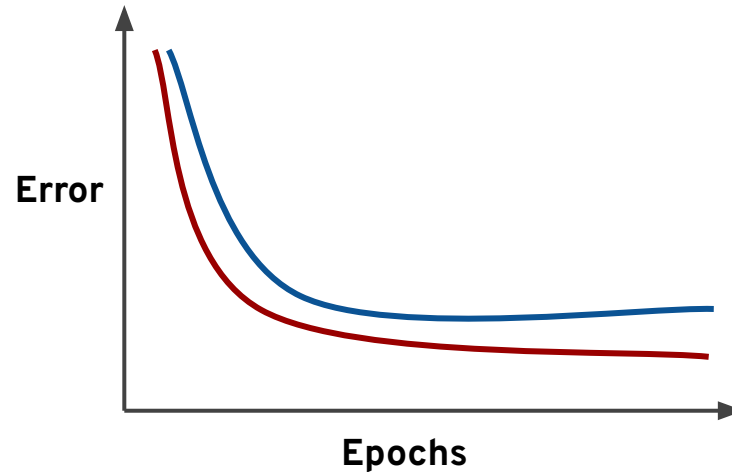
- Next we check performance on the **test set**





Machine Learning

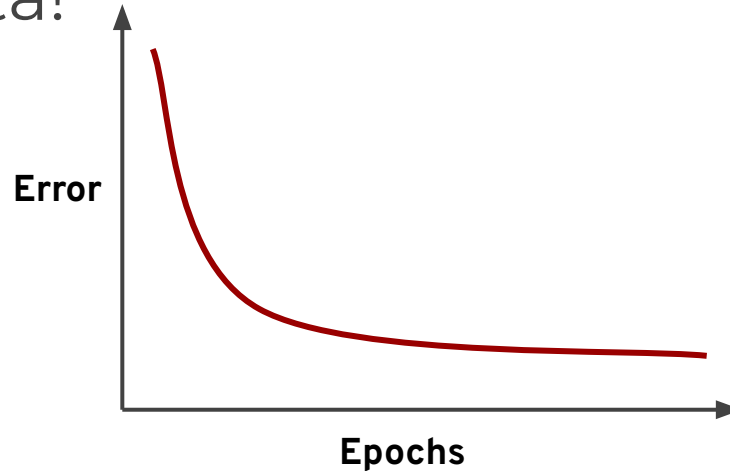
- Ideally the model would perform well on both, with similar behavior.





Machine Learning

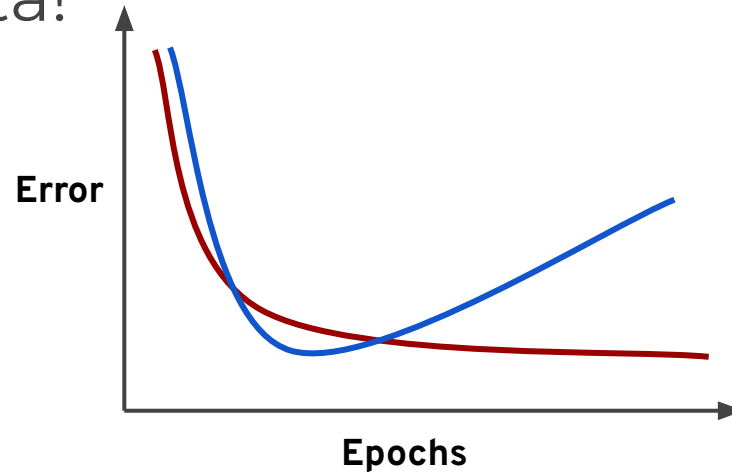
- But what happens if we overfit on the training data? That means we would perform poorly on new test data!





Machine Learning


- But what happens if we overfit on the training data? That means we would perform poorly on new test data!





Machine Learning

- This is a good indication of training too much on the training data, you should look for the point to cut off training time!





Machine Learning

- We'll check on this idea again when we actually begin creating models!
- For now just be aware of this possible issue!



Evaluating Performance

CLASSIFICATION



Model Evaluation

- We just learned that after our machine learning process is complete, we will use performance metrics to evaluate how our model did.
- Let's discuss classification metrics in more detail!



Model Evaluation

- The key classification metrics we need to understand are:
 - Accuracy
 - Recall
 - Precision
 - F1-Score



Model Evaluation

- But first, we should understand the reasoning behind these metrics and how they will actually work in the real world!



Model Evaluation

- Typically in any classification task your model can only achieve two results:
 - Either your model was **correct** in its prediction.
 - Or your model was **incorrect** in its prediction.



Model Evaluation

- Fortunately incorrect vs correct expands to situations where you have multiple classes.
- For the purposes of explaining the metrics, let's imagine a **binary classification** situation, where we only have two available classes.



Model Evaluation

- In our example, we will attempt to predict if an image is a dog or a cat.
- Since this is supervised learning, we will first **fit/train** a model on **training data**, then **test** the model on **testing data**.
- Once we have the model's predictions from the **X_test** data, we compare it to the **true y values** (the correct labels).



Model Evaluation

TRAINED
MODEL



Model Evaluation



**Test Image
from X_{test}**

TRAINED
MODEL



Model Evaluation



**Test Image
from X_{test}**

DOG

**Correct Label
from y_{test}**





Model Evaluation



**Test Image
from X_{test}**

DOG

**Correct Label
from y_{test}**

TRAINED
MODEL

DOG

**Prediction on
Test Image**



Model Evaluation



**Test Image
from X_{test}**


**TRAINED
MODEL**

DOG

**Prediction on
Test Image**

DOG

**Correct Label
from y_{test}**

 **DOG == DOG ?**

Compare Prediction to Correct Label



Model Evaluation



**Test Image
from X_{test}**

**TRAINED
MODEL**

CAT

**Prediction on
Test Image**

DOG

**Correct Label
from y_{test}**



$DOG == CAT ?$

Compare Prediction to Correct Label



Model Evaluation

- We repeat this process for all the images in our X test data.
- At the end we will have a count of correct matches and a count of incorrect matches.
- The key realization we need to make, is that **in the real world, not all incorrect or correct matches hold equal value!**



Model Evaluation

- Also in the real world, a single metric won't tell the complete story!
- To understand all of this, let's bring back the 4 metrics we mentioned and see how they are calculated.
- We could organize our predicted values compared to the real values in a **confusion matrix**.



Model Evaluation

- Accuracy
 - Accuracy in classification problems is the **number of correct predictions** made by the model divided by the **total number of predictions**.



Model Evaluation

- Accuracy
 - For example, if the X_test set was 100 images and our model **correctly** predicted 80 images, then we have **80/100**.
 - **0.8** or **80% accuracy**.



Model Evaluation

- Accuracy
 - Accuracy is useful when target classes are well balanced
 - In our example, we would have roughly the same amount of cat images as we have dog images.



Model Evaluation

- Accuracy
 - Accuracy is **not** a good choice with **unbalanced** classes!
 - Imagine we had 99 images of dogs and 1 image of a cat.
 - If our model was simply a line that always predicted **dog** we would get 99% accuracy!



Model Evaluation

- Accuracy
 - Imagine we had 99 images of dogs and 1 image of a cat.
 - If our model was simply a line that always predicted **dog** we would get 99% accuracy!
 - In this situation we'll want to understand **recall** and **precision**



Model Evaluation

- Recall
 - Ability of a model to find all the relevant cases within a dataset.
 - The precise definition of recall is the number of true positives **divided by** the number of true positives plus the number of false negatives.



Model Evaluation

- Precision
 - Ability of a classification model to identify only the relevant data points.
 - Precision is defined as the number of true positives divided by the number of true positives plus the number of false positives.



Model Evaluation

- Recall and Precision
 - Often you have a trade-off between Recall and Precision.
 - While recall expresses the ability to find all relevant instances in a dataset, precision expresses the proportion of the data points our model says was relevant actually were relevant.



Model Evaluation

- F1-Score
 - In cases where we want to find an optimal blend of precision and recall we can combine the two metrics using what is called the F1 score.



Model Evaluation

- F1-Score
 - The F1 score is the harmonic mean of precision and recall taking both metrics into account in the following equation:

$$F_1 = 2 * \frac{\textit{precision} * \textit{recall}}{\textit{precision} + \textit{recall}}$$



Model Evaluation

- F1-Score
 - We use the harmonic mean instead of a simple average because it punishes extreme values.
 - A classifier with a precision of 1.0 and a recall of 0.0 has a simple average of 0.5 but an F1 score of 0.



Model Evaluation

- We can also view all correctly classified versus incorrectly classified images in the form of a confusion matrix.



Confusion Matrix

		predicted condition	
		prediction positive	prediction negative
true condition	condition positive	True Positive (TP)	False Negative (FN) (type II error)
	condition negative	False Positive (FP) (Type I error)	True Negative (TN)



Confusion Matrix

		predicted condition		
total population		prediction positive	prediction negative	Prevalence $= \frac{\Sigma \text{ condition positive}}{\Sigma \text{ total population}}$
true condition	condition positive	True Positive (TP)	False Negative (FN) (type II error)	True Positive Rate (TPR), Sensitivity, Recall, Probability of Detection $= \frac{\Sigma \text{ TP}}{\Sigma \text{ condition positive}}$
	condition negative	False Positive (FP) (Type I error)	True Negative (TN)	False Positive Rate (FPR), Fall-out, Probability of False Alarm $= \frac{\Sigma \text{ FP}}{\Sigma \text{ condition negative}}$
Accuracy $= \frac{\Sigma \text{ TP} + \Sigma \text{ TN}}{\Sigma \text{ total population}}$		Positive Predictive Value (PPV), Precision $= \frac{\Sigma \text{ TP}}{\Sigma \text{ prediction positive}}$	False Omission Rate (FOR) $= \frac{\Sigma \text{ FN}}{\Sigma \text{ prediction negative}}$	Positive Likelihood Ratio (LR+) $= \frac{\text{TPR}}{\text{FPR}}$
		False Discovery Rate (FDR) $= \frac{\Sigma \text{ FP}}{\Sigma \text{ prediction positive}}$	Negative Predictive Value (NPV) $= \frac{\Sigma \text{ TN}}{\Sigma \text{ prediction negative}}$	Negative Likelihood Ratio (LR-) $= \frac{\text{FNR}}{\text{TNR}}$



Model Evaluation

- The main point to remember with the confusion matrix and the various calculated metrics is that they are all fundamentally ways of comparing the predicted values versus the true values.
- What constitutes “good” metrics, will really depend on the specific situation!



Model Evaluation

- Still confused on the confusion matrix?
- No problem! Check out the Wikipedia page for it, it has a really good diagram with all the formulas for all the metrics.
- Throughout the training, we'll usually just print out metrics (e.g. accuracy).



Model Evaluation

- Let's think back on this idea of:
 - What is a good enough accuracy?
- This all depends on the context of the situation!
- Did you create a model to predict presence of a disease?
- Is the disease presence well balanced in the general population? (Probably not!)



Model Evaluation

- Often models are used as quick diagnostic tests to have **before** having a more invasive test (e.g. getting urine test before getting a biopsy)
- We also need to consider what is at stake!



Model Evaluation

- Often we have a precision/recall trade off, We need to decide if the model will should focus on fixing False Positives vs. False Negatives.
- In disease diagnosis, it is probably better to go in the direction of False positives, so we make sure we correctly classify as many cases of disease as possible!



Model Evaluation

- All of this is to say, machine learning is not performed in a “vacuum”, but instead a collaborative process where we should consult with experts in the domain (e.g. medical doctors)



Evaluating Performance

REGRESSION



Evaluating Regression

- Let's take a moment now to discuss evaluating Regression Models
- Regression is a task when a model attempts to predict continuous values (unlike categorical values, which is classification)



Evaluating Regression

- You may have heard of some evaluation metrics like accuracy or recall.
- These sort of metrics aren't useful for regression problems, we need metrics designed for **continuous** values!



Evaluating Regression

- For example, attempting to predict the price of a house given its features is a **regression task**.
- Attempting to predict the country a house is in given its features would be a classification task.



Evaluating Regression

- Let's discuss some of the most common evaluation metrics for regression:
 - Mean Absolute Error
 - Mean Squared Error
 - Root Mean Square Error



Evaluating Regression

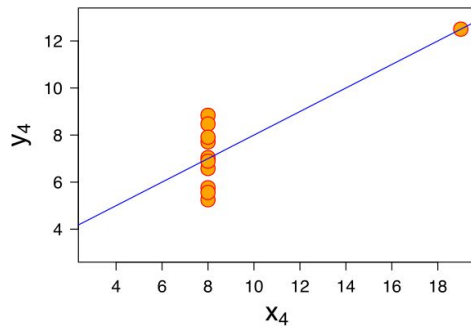
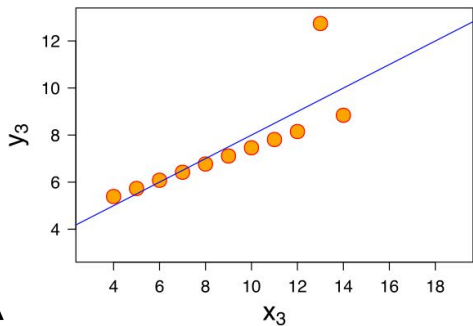
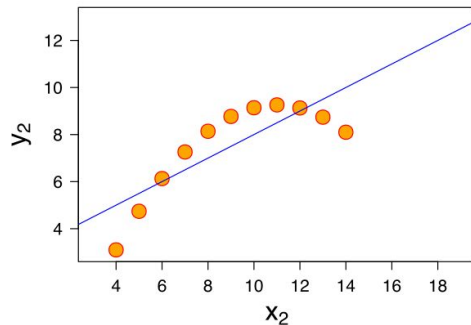
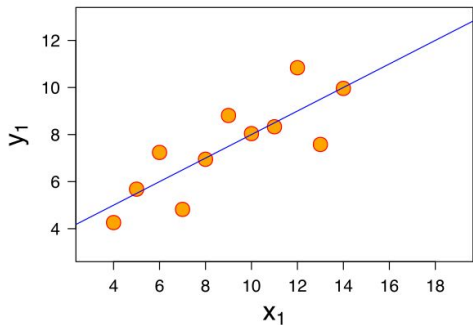
- Mean Absolute Error (MAE)
 - This is the mean of the absolute value of errors.
 - Easy to understand

$$\frac{1}{n} \sum_{i=1}^n |y_i - \hat{y}_i|$$



Evaluating Regression


- MAE won't punish large errors however.





Evaluating Regression

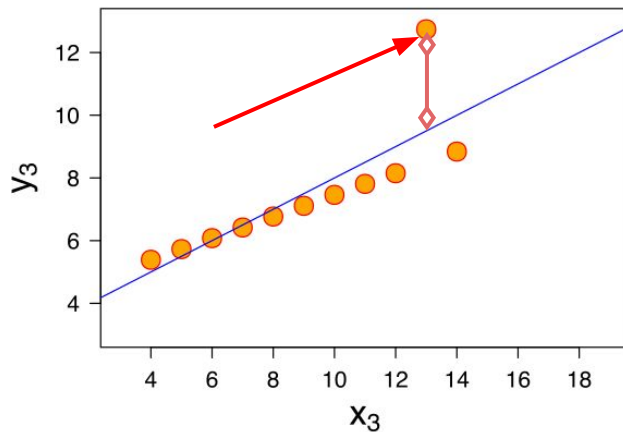
- MAE won't punish large errors however.





Evaluating Regression

- We want our error metrics to account for these!





Evaluating Regression

- Mean Squared Error (MSE)
 - This is the mean of the squared errors.
 - Larger errors are noted more than with MAE, making MSE more popular.

$$\frac{1}{n} \sum_{i=1}^n (y_i - \hat{y}_i)^2$$



Evaluating Regression

- Root Mean Square Error (RMSE)
 - This is the root of the mean of the squared errors.
 - Most popular (has same units as y)

$$\sqrt{\frac{1}{n} \sum_{i=1}^n (y_i - \hat{y}_i)^2}$$



Machine Learning

- Most common question from students:
 - “Is this value of RMSE good?”
- Context is everything!
- A RMSE of \$10 is fantastic for predicting the price of a house, but horrible for predicting the price of a candy bar!



Machine Learning

- Compare your error metric to the average value of the label in your data set to try to get an intuition of its overall performance.
- Domain knowledge also plays an important role here!



Machine Learning

- Context of importance is also necessary to consider.
- We may create a model to predict how much medication to give, in which case small fluctuations in RMSE may actually be very significant.



Evaluating Regression

- You should now feel comfortable with the various methods of evaluating a regression task.



Unsupervised Learning



Machine Learning

- We've covered supervised learning, where the **label was known** due to **historical labeled data**.
- But what happens when we don't have historical labels?



Machine Learning

- There are certain tasks that fall under unsupervised learning:
 - Clustering
 - Anomaly Detection
 - Dimensionality Reduction



Machine Learning

- Clustering
 - Grouping together **unlabeled** data points into categories/clusters
 - Data points are assigned to a cluster based on similarity



Machine Learning

- Anomaly Detection
 - Attempts to detect outliers in a dataset
 - For example, fraudulent transactions on a credit card.



Machine Learning

- Dimensionality Reduction
 - Data processing techniques that reduces the number of features in a data set, either for compression, or to better understand underlying trends within a data set.



Machine Learning

- Unsupervised Learning
 - It's important to note, these are situations where we **don't** have the correct answer for historical data!
 - Which means evaluation is much harder and more nuanced!



Unsupervised Process





Machine Learning

- Later on in the course, we'll explore unsupervised learning processes with specialized neural network structures, such as autoencoders.